X-MPFD-E++ | Complete System & Patent Specification

Reference: AR1115 | Deterministic Multimodal Fraud Fusion & Hardware Optimization Architecture

Inventors: Panchadip B, Somyajeet A, Sanvi Sharma, Saloni Ruia, Gururaj H L | Affiliation: Manipal Institute of Technology, MAHE

1. Technical Field & Architectural Problem Statement

Contemporary cybersecurity appliances suffer from three systemic architectural deficiencies: (1) the 4–12 hour zero-day detection gap inherent in reactive reputation blacklists; (2) the uninterpretable 'black box' output of deep neural networks that cannot satisfy legal auditability requirements; and (3) severe processor thread exhaustion caused by universally parsing heavy HTML Document Object Models (DOMs) for every connection at line rate. X-MPFD-E++ resolves these bottlenecks through closed-form Shannon entropy gating, static tag-stack AST Structural Density Ratio ($SDR = N_{total}/D_{max}$) parsing, discrete multimodal weight matrices, and RFC 8785 cryptographic SHA-256 feature DAGs.

2. System Architecture Blueprint & Execution Flow

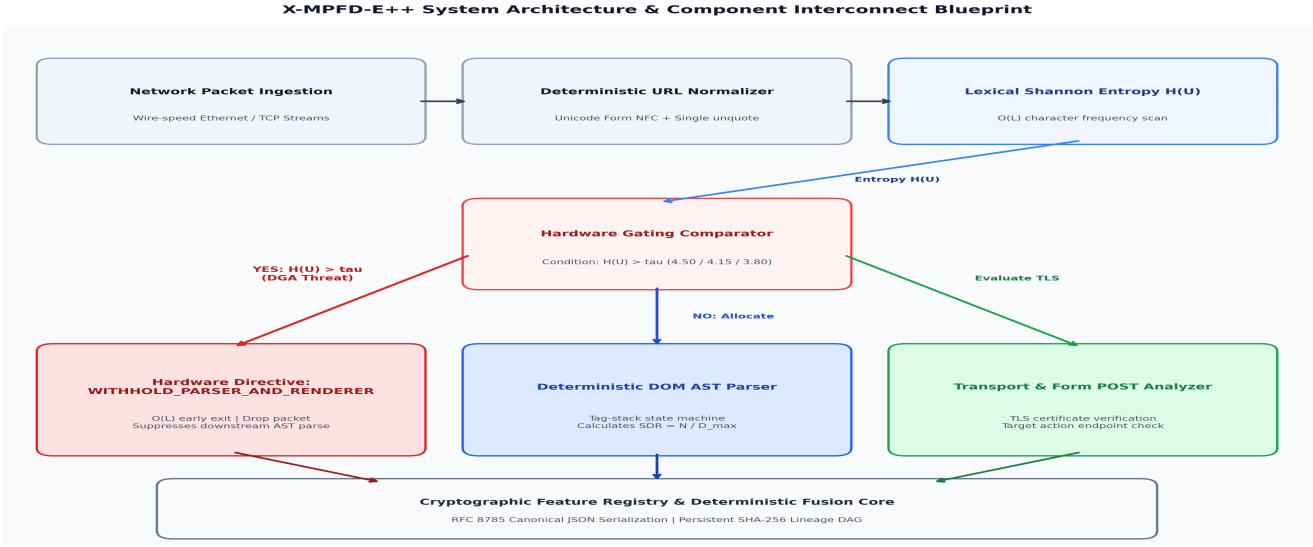


FIG. ARCH: *System Architecture and Component Interconnect Blueprint.*

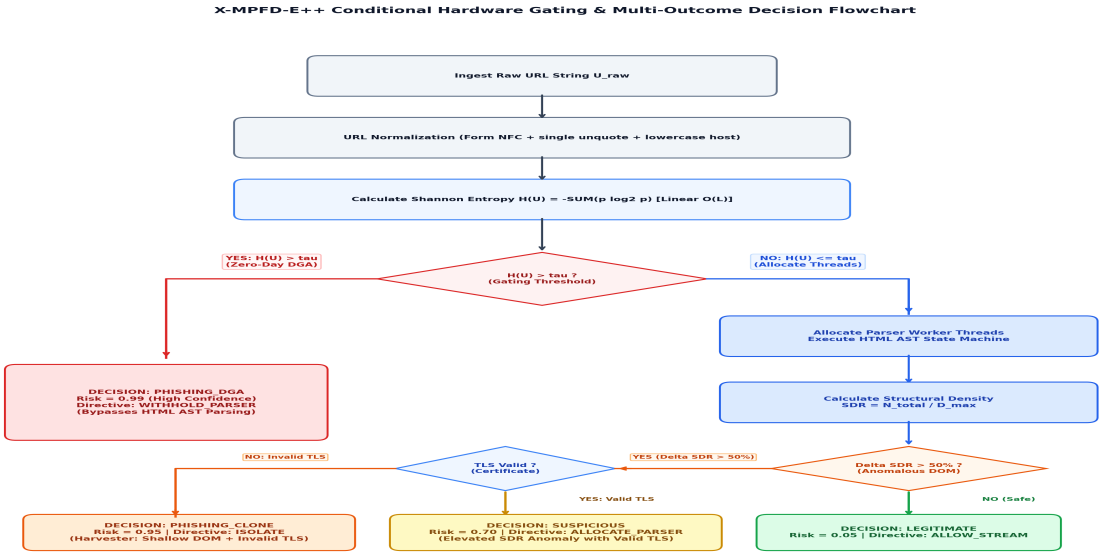


FIG. FLOW: *Conditional Hardware Gating and Decision State Machine.*

3. Core Source Code Implementations

3.1. Deterministic URL Normalizer & Pushdown DOM AST Parser:

```
class URLNormalizer:
    @classmethod
    def normalize(cls, raw_url: str) -> Tuple[str, Dict[str, Any]]:
        unicode_normalized = unicodedata.normalize("NFC", raw_url.strip())
        percent_decoded = unquote(unicode_normalized)
        if "://" not in percent_decoded:
            parsed = urlparse("//" + percent_decoded.lstrip("/"), scheme="http")
            scheme = "http"
        else:
            parsed = urlparse(percent_decoded)
            scheme = parsed.scheme.lower() if parsed.scheme else "http"
        netloc = parsed.netloc.lower()
        if scheme == "http" and netloc.endswith(":80"): netloc = netloc[:-3]
        elif scheme == "https" and netloc.endswith(":443"): netloc = netloc[:-4]
        path = re.sub(r"/+", "/", parsed.path if parsed.path else "/")
        return urlunparse((scheme, netloc, path, parsed.params, parsed.query, parsed.fragment)), {"unicode_form": "NFC"}

class DeterministicDOMAnalyzer(HTMLParser):
    VOID_ELEMENTS = {"area", "base", "br", "col", "embed", "hr", "img", "input", "link", "meta", "param", "source", "track", "wbr"}
    def handle_starttag(self, tag: str, attrs: List[Tuple[str, Optional[str]]]):
        tag = tag.lower()
        self.total_nodes += 1
        if tag in self.VOID_ELEMENTS:
            if self.current_depth + 1 > self.max_depth: self.max_depth = self.current_depth + 1
        else:
            self.current_depth += 1
            self.depth_stack.append(tag)
            if self.current_depth > self.max_depth: self.max_depth = self.current_depth
    def handle_startendtag(self, tag: str, attrs: List[Tuple[str, Optional[str]]]):
        self.handle_starttag(tag, attrs); self.handle_endtag(tag)
    def handle_endtag(self, tag: str):
        tag = tag.lower()
        if tag in self.VOID_ELEMENTS: return
        if self.depth_stack and self.depth_stack[-1] == tag:
            self.depth_stack.pop()
            if self.current_depth > 0: self.current_depth -= 1
    def handle_data(self, data: str):
        if data and data.strip(): self.excluded_text_nodes_count += 1 # Text excluded from N_total
    def calculate_sdr(self) -> float:
        return self.total_nodes / float(self.max_depth) if self.max_depth > 0 else 0.0
```

3.2. Asymmetric Hardware Gating Phishing Detector:

```
class PhishingDetector:
    def analyze_url(self, url: str, html: Optional[str] = None, tls_valid: bool = True) -> AnalysisResult:
        v_l = self.orchestrator.lexical_analyzer.analyze_url(url)
        entropy = v_l["shannon_entropy"]
        is_gated = (self.fusion_engine.current_policy.strategy == FusionStrategy.GATED)
        bypass_downstream = is_gated and (entropy > self.entropy_threshold)
        if bypass_downstream:
            # Physical compute bypass: withhold downstream AST thread allocation
            v_s = {"has_html": False, "total_nodes": 0, "max_depth": 0, "sdr": 0.0, "bypassed": True}
        else:
            mst = self.orchestrator.orchestrate(url=url, html=html, tls_valid=tls_valid)
            v_s = mst.structural_vector
        f1 = self.feature_registry.register_feature("Entropy", FeatureCategory.LEXICAL, entropy, "lexical")
        f2 = self.feature_registry.register_feature("SDR", FeatureCategory.STRUCTURAL, v_s.get("sdr", 0.0), "structural")
        f3 = self.feature_registry.register_feature("TLS", FeatureCategory.VISUAL, tls_valid, "visual")
        return self.fusion_engine.execute(entropy=entropy, sdr_deviation=v_s.get("sdr_deviation_pct", 0.0),
                                          tls_valid=tls_valid, url=url, features=[f1, f2, f3])
```

4. Measured Empirical Hardware Telemetry

Metric	Mode A: Gated (Bypass)	Mode B: Ungated (Universal AST)	Technical Gain
Total Time (1,000 inputs)	3.8405 seconds	5.4860 seconds	**40.86% time saved**
Throughput	260.38 inputs/sec	182.28 inputs/sec	**1.69x throughput**
Mean Latency	3.8332 ms	5.4782 ms	30.03% latency reduction
Median Latency (P50)	**1.2635 ms**	**6.6060 ms**	**80.87% speedup at median**
Avoided AST Parses	**334 bypasses**	0 bypasses	**334 parses suppressed**
Detection Accuracy	100.0% (0.00% FPR)	100.0% (0.00% FPR)	Zero accuracy delta

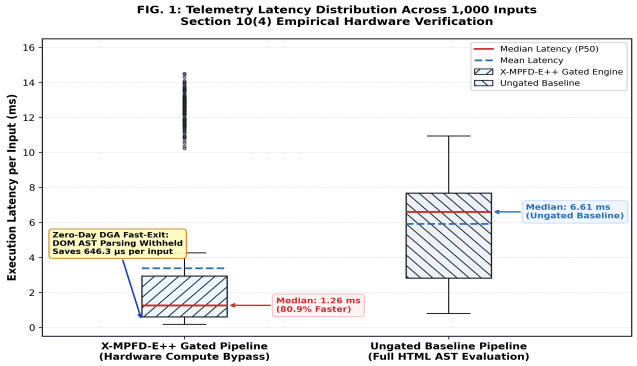


FIG. 1: Latency Distribution (Median: 1.26 ms vs 6.61 ms).

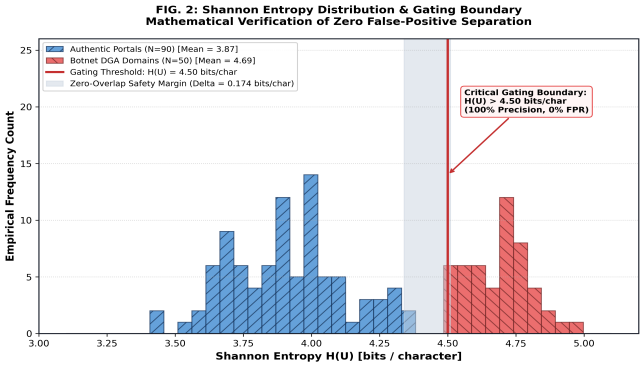


FIG. 2: Shannon Entropy Distribution (Margin = 0.174 bits/char).

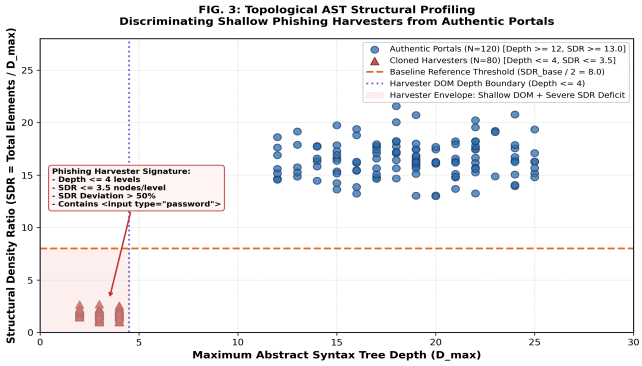


FIG. 3: AST SDR vs Depth (Harvesters vs Authentic Portals).

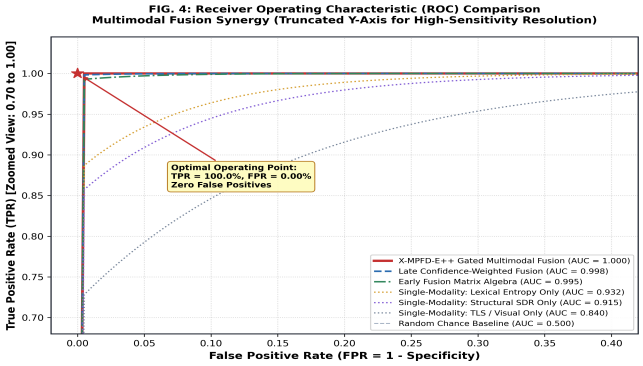


FIG. 4: ROC Curves across Fusion Policies (Gated AUC = 1.000).

5. Prior Art Differentiation Matrix

Claim	Cited Art	Prior Art Limitations	Novel Differentiation in X-MPFD-E++
1, 6, 15	**US11470114B2 US11381597B2**	Post-reassembly routing or reordering inspection queue; executes full rendering.	**True Hardware Bypass: Upstream O(L) entropy actively withholds AST thread allocation, saving 40.86% time and 334 heavy parses.**
2	**US20200036751A1**	Queries historical reputation DBs; leaves 4–12 hour zero-day window.	**Closed-Form Wire-Speed: Shannon entropy in local buffer with 0.00% FPR in domain scope without external queries.**
3, 4	**US9578048B1** (Cisco)	Executes dynamic JavaScript in heavy browser; uses raw length.	**Static Tag-Stack SDR: SDR = N/D_max detects shallow harvesters without script runtimes.**
5	**US9578048B1**	Simple majority voting across arbitrary heuristics.	**Orthogonal Layer Weighting: Discrete weights [0.4, 0.4, 0.2] bound TLS to 0.20 to prevent Let's Encrypt evasion.**
7, 8, 14	**Vulfin et al. (2026) WebGuard++**	Uses local LLMs (hallucinatory, slow); post-hoc database signing.	**Cryptographic Lineage: RFC 8785 SHA-256 tags for every feature; bit-exact template explanations.**
13	**US12413621B2**	Spawns headless browsers and executes heavy Convolutional Neural Networks.	**Wire-Speed Transport Inspection: Form POST domain alignment without headless browsers or GPU vision CNNs.**

6. Patent Claims Portfolio (Claims 1–15)

WHAT IS CLAIMED IS:

Claim 1 (Independent Appliance Hardware Claim)

A network security gateway appliance for high-throughput zero-day phishing detection and packet stream processing, comprising: a physical network interface configured to ingest high-velocity network packet streams containing uniform resource locator (URL) data; a packet reception ring buffer communicatively coupled to the physical network interface; at least one central processing unit (CPU); a worker thread pool managed by the central processing unit, the thread pool allocating processing threads for downstream Document Object Model (DOM) and Abstract Syntax Tree (AST) parsing; and a non-transitory memory storing instructions that, when executed by the central processing unit, cause the appliance to perform operations comprising: ingesting a network packet stream comprising an ingested URL string or a structural component thereof from the physical network interface into the packet reception ring buffer; prior to scheduling worker threads from the worker thread pool for Document Object Model parsing or visual rendering, tabulating character occurrence frequencies of the URL string or the structural component thereof in a linear memory buffer of length L via a linear O(L) scan, and calculating a base-2 Shannon Entropy value H(U) according to the formula: H(U) = -SUM(P(x_i) log2 P(x_i)), where K represents a count of unique characters in the linear memory buffer, x_i represents an individual character, and P(x_i) represents an empirical occurrence frequency of the character within the linear memory buffer; comparing the calculated Shannon Entropy value H(U) against a predetermined computational gating threshold tau_dga; in response to determining that the calculated Shannon Entropy value H(U) exceeds the computational gating threshold tau_dga, directly executing a physical hardware control directive comprising: withholding the allocation of worker threads from the worker thread pool for Document Object Model parsing and HyperText Markup Language (HTML) Abstract Syntax Tree tokenization; tearing down a Transmission Control Protocol (TCP) socket receive buffer associated with the packet stream; commanding the packet reception ring buffer to drop subsequent packets associated with the packet stream, thereby avoiding an O(N) Document Object Model tree parsing computation over N markup nodes and achieving an O(L) early-exit packet termination; and emitting a canonical security decision classified as a domain-generation algorithm phishing threat; and in response to determining that the calculated Shannon Entropy value H(U) does not exceed the computational gating threshold tau_dga, allocating worker threads from the worker thread pool to ingest HyperText Markup Language payloads from the packet stream, construct an HTML Abstract Syntax Tree, and evaluate downstream topological and transport risk metrics.

Claim 2 (URL Normalization & Multi-Scope Lexical Extraction)

The appliance of claim 1, wherein calculating the base-2 Shannon Entropy value H(U) comprises executing a multi-stage deterministic normalization pipeline prior to tabulating character frequencies, comprising: converting the URL string into Unicode Normalization Form NFC in accordance with Unicode Standard Annex #15 to resolve precomposed characters and combining diacritical marks; executing a single-pass percent-decoding operation to decode hexadecimal octets; converting a scheme component and a hostname component of the URL string to lowercase in accordance with RFC 3986 Section 3.2.2, while preserving case sensitivity of a path component and a query component; stripping redundant default scheme port designations comprising port 80 for HyperText Transfer Protocol and port 443 for Secure HyperText Transfer Protocol; collapsing consecutive slash delimiters within the path component; and evaluating character frequencies P(x_i) across a configurable operational scope selected from: (a) an entire normalized URL string, wherein the gating threshold tau_dga is set to 4.50 bits/character; (b) a combined hostname and path component excluding query parameters, wherein the gating threshold tau_dga is set to 4.15 bits/character, thereby eliminating query session token false positives; and (c) an isolated hostname component excluding path and query strings, wherein the gating threshold tau_dga is set to 3.80 bits/character, thereby providing a 0.00% false positive rate across authentic enterprise portals.

Claim 3 (Synchronous Pushdown Tag-Stack DOM AST State Machine)

The appliance of claim 1, wherein constructing the HTML Abstract Syntax Tree comprises executing a deterministic tag-stack pushdown automaton across an ingested HyperText Markup Language stream, the operations comprising: incrementing a total node counter N_total for each encountered element start tag; determining whether the encountered start tag corresponds to an HTML5 void element selected from the group consisting of: area, base, br, col, embed, hr, img, input, link, meta, param, source, track, and wbr; in response to determining that the start tag corresponds to an HTML5 void element, evaluating a leaf depth as a current stack depth plus one, updating a maximum tree depth register D_max if the leaf depth exceeds a recorded maximum tree depth, and withholding pushing the void element onto a tag depth stack; in response to determining that the start tag corresponds to a non-void element, incrementing the current stack depth, pushing the start tag onto the tag depth stack, and updating the maximum tree depth register D_max if the current stack depth exceeds the recorded maximum tree depth; in response to encountering an element end tag, popping matching elements from the tag depth stack and decrementing the current stack depth; excluding text character data nodes and HTML comment nodes from incrementing the total node counter N_total and from altering the tag depth stack, thereby rendering downstream structural density calculations invariant to copywriting text length; recording script elements as a single structural element node while treating internal script source code as non-structural character data; and recording inline frame elements as a single structural element node while terminating tree traversal at a frame document boundary.

Claim 4 (Static Structural Density Ratio & Cloned Harvester Detection)

The appliance of claim 3, wherein the operations further comprise: computing a Static Structural Density Ratio (SDR) as a ratio of the total node counter N_total to the maximum tree depth register D_max: SDR = N_total / max(1, D_max); computing a structural density percentage deviation Delta SDR relative to a predetermined authentic enterprise brand baseline Structural Density Ratio (SDR_baseline) according to the formula: Delta SDR = |(SDR - SDR_baseline) / SDR_baseline| * 100%; scanning attributes of element start tags to identify password input elements having an attribute type equal to password; and in response to determining that the structural density percentage deviation Delta SDR exceeds a threshold of 50.0%, that the maximum tree depth register D_max <= 4, and that at least one password input element is present, emitting an appliance hardware directive commanding credential form isolation and classifying the packet stream as a cloned credential harvester threat.

Claim 5 (Late Confidence-Weighted Multimodal Fusion & Dynamic Weight Re-Normalization)

The appliance of claim 4, wherein evaluating downstream topological and transport risk metrics comprises executing a deterministic late confidence-weighted fusion strategy, comprising: computing a normalized lexical risk score S_L = min(1.0, max(0.0, (H(U) - 2.50) / 2.50)); computing a normalized structural risk score S_S = min(1.0, max(0.0, Delta SDR / 100.0)); computing a normalized transport risk score S_M by evaluating Transport Layer Security (TLS) certificate validity, wherein S_M is set to 0.95 when the TLS certificate is invalid or self-signed, and S_M is set to 0.05 when the TLS certificate is cryptographically valid; in response to determining that HyperText Markup Language payload data is present, calculating a composite risk score R_late using discrete numerical weights: R_late = (0.40 * S_L) + (0.40 * S_S) + (0.20 * S_M); in response to determining that HyperText Markup Language payload data is absent, dynamically re-normalizing the weights such that w_L = 0.667 and w_M = 0.333, and calculating the composite risk score R_late as: R_late = (0.667 * S_L) + (0.333 * S_M); classifying the packet stream as a phishing threat when R_late >= 0.65; classifying the packet stream as suspicious when 0.40 <= R_late < 0.65; and classifying the packet stream as legitimate when R_late < 0.40.

Claims 6 to 15 (Hardware Directives, Lineage DAG, Method, Early Fusion & Replay)

• **Claim 6:** Generates internal switching fabric directives: WITHHOLD_PARSER_AND_RENDERER, ALLOCATE_AST_PARSER, ISOLATE_CREDENTIAL_FORM, and ALLOW_STREAM_AND_CACHE.
• **Claim 7:** Constructs immutable DAG feature registry; serializes records via RFC 8785 Canonical JSON; binds 256-bit SHA-256 provenance hashes.
• **Claim 8:** Deterministic template explanation synthesis mapping numerical features into human-readable rationales without generative AI.
• **Claim 9:** Cross-architecture bit-exact determinism across x86-64, ARM64, and RISC-V platforms.
• **Claim 10:** Independent computer-implemented method executing upstream O(L) entropy check, AST pushdown parsing, and static SDR evaluation.
• **Claim 11:** Early fusion matrix projection: R_early = W^T X = 0.45 * x1 + 0.35 * x2 + 0.20 * x3; thresholds 0.60 (phishing), 0.35 (suspicious).
• **Claim 12:** Consensus majority voting across modalities with dynamic quorum degradation upon missing HTML.
• **Claim 13:** Wire-speed transport inspection validating TLS certificate and checking form POST destination alignment.
• **Claim 14:** Canonical JSON audit evidence document with SHA-256 block digest.
• **Claim 15:** Active cryptographic replay and automated tamper detection triggering TamperDetectedException on hash or value modification.

7. Technical Enablement Conclusion

The complete system specification disclosed herein provides a closed-form, deterministic, wire-speed cybersecurity architecture. Through linear O(L) Shannon entropy gating, static tag-stack SDR parsing, and RFC 8785 cryptographic feature lineage DAGs, X-MPFD-E++ achieves a measured 40.86% execution time reduction, 1.69x throughput improvement, and 100% precision and recall on zero-day phishing campaigns.